Transform then pool or pool and then transform?

Luiz Max F. de Carvalho

January 3, 2019

Abstract

In this note I discuss the order of operations between transforming and pooling a set of distributions in relation to whether the transform in question is invertible. I give an explicit example where both procedures lead to the same density even when the transform is not invertible and also discuss an example where the transform-then-pool and pool-then-transform with a non-invertible transform lead to distributions in the same family which are nonetheless distinct. The lesson to be learned is that while the order of operations usually matters, it may not, depending on the distributions under consideration.

Key-words: logarithmic pooling; invertible; log-normal.

Background

Logarithmic pooling is a popular method for combining opinions on an agreed quantity, specially when these opinions can be framed as probability distributions. Let $\mathbf{F}_\theta := \{f_0(\theta), f_1(\theta), \ldots, f_K(\theta)\}$ be a set of distributions representing the opinions of $K + 1$ experts and let $\boldsymbol{\alpha} := \{\alpha_0, \alpha_1, \ldots, \alpha_K\} \in \mathcal{S}^K$ be the vector of weights, such that $\alpha_i > 0 \, \forall i$ and $\sum_{i=0}^{K} \alpha_i = 1$, i.e., $\mathcal{S}^{K+1}$ is the space of all open simplices of dimension $K + 1$. The **logarithmic pooling operator** $\mathcal{LP}(\mathbf{F}_\theta, \boldsymbol{\alpha})$ is defined as

$$\mathcal{LP}(\mathbf{F}_\theta, \boldsymbol{\alpha}) := \pi(\theta|\boldsymbol{\alpha}) = t(\boldsymbol{\alpha}) \prod_{i=0}^{K} f_i(\theta)^{\alpha_i}, \tag{1}$$

where $t(\boldsymbol{\alpha}) = \int_{\boldsymbol{\Theta}} \prod_{i=0}^{K} f_i(\theta)^{\alpha_i} d\theta$. This pooling method enjoys several desirable properties and yields tractable distributions for a large class of distribution families (Genest et al., 1984).

Pool then transform or transform and then pool?

Definition 1. *Let $A, B \subseteq \mathbb{R}^p$. A function $h : A \to B$ is **invertible** iff $\exists \, h^{-1} : B \to A$ with $h^{-1}(h(a)) = a \, \forall \, a \in A$. Let π_A be an arbitrary probability measure in A. If h is monotonic and differentiable we can write $\pi_B(B) = \pi_A(h^{-1}(A))|J|$, where $|J|$ is the absolute determinant of the Jacobian matrix with entries $J_{ik} := \partial h_k^{-1}/\partial a_i$, $i, k = 1, 2, \ldots, p$.*

Definition 2. *Let f and g be probability density functions with support on $\mathcal{X} \subseteq \mathbb{R}^d$. We say $f \equiv g$ if $f(x) = g(x)$ for almost every point $x \in \mathcal{X}$, that is, almost everywhere, a.e. This equivalent to stating that $\mu(A) = 0$ where $A := \{x \in \mathcal{X} : f(x) \neq g(x)\}$ and μ is the dominating measure.*

Suppose the we are interested in the distribution of a random variable $Y \in \mathcal{Y} \subseteq \mathbb{R}^q$ when one has a random variable $X \in \mathcal{X} \subseteq \mathbb{R}^p$ with $\phi : \mathcal{X} \to \mathcal{Y}$. Let $|J_\phi|$ be the (absolute) Jacobian determinant w.r.t. ϕ. Suppose further that each expert i produces a distribution $f_i(X)$ such that we can construct the object $\mathbf{F}_X = \{f_1(X), f_2(X), \ldots, f_K(X)\}$. Then one can either:

(a) **Pool-then-transform:** construct $\pi_X(X|\boldsymbol{\alpha}) = \mathcal{LP}(\mathbf{F}_X, \boldsymbol{\alpha})$ and then apply ϕ to obtain $\pi_Y(Y|\boldsymbol{\alpha}) := \pi_X(\phi^{-1}(Y)|\boldsymbol{\alpha})|J_\phi|$;

(b) **Transform-then-pool:** apply the transform to each component i of $\mathbf{F}_X$ to build

$$\mathbf{G}_Y := \{g_i(Y), g_2(Y), \ldots, g_K(Y)\}$$

and obtain $\pi'_Y(Y|\boldsymbol{\alpha}) = \mathcal{LP}(\mathbf{G}_Y, \boldsymbol{\alpha})$.

Remark 1. *If ϕ is invertible, then $\pi_Y(Y|\boldsymbol{\alpha}) \equiv \pi'_Y(Y|\boldsymbol{\alpha})$.*

Proof. First,

$$\pi_Y(y|\boldsymbol{\alpha}) \propto \pi_X(\phi^{-1}(y))|J_\phi|, \tag{2}$$

$$= \prod_{i=0}^{K} \left[f_i(\phi^{-1}(y)) \right]^{\alpha_i} |J_\phi|. \tag{3}$$

For situation (b) we have:

$$g_i(y) = f_i(\phi^{-1}(y))|J_\phi|. \tag{4}$$

And,

$$\pi'_Y(y|\boldsymbol{\alpha}) \propto \prod_{i=0}^{K} g_i(y)^{\alpha_i} \tag{5}$$

$$= \prod_{i=0}^{K} \left[f_i(\phi_x^{-1}(y))|J_\phi| \right]^{\alpha_i} \tag{6}$$

$$= \prod_{i=0}^{K} \left[f_i(\phi^{-1}(y)) \right]^{\alpha_i} |J_\phi|, \tag{7}$$

as claimed. $\square$

An interesting idea is whether Remark 1 is an iff result.

Remark 2. *It is possible to have $\pi_Y(\cdot|\boldsymbol{\alpha}) \equiv \pi'_Y(\cdot|\boldsymbol{\alpha})$ even when ϕ is not invertible.*

Proof. We will show this by way of an explicit example. In general, we can define

$$g_i(y) = \sum_{\omega \in \Omega(y)} \frac{f_i(\omega)}{|\phi'(\omega)|}. \tag{8}$$

where $\Omega(y) := \{x : \phi(x) = y\}$ [1]. Notice that for any $x \in \mathcal{X}$, there exists y such that $x \in \Omega(y)$. Assume that the transform ϕ is non-invertible, i.e., that $|\Omega(y)| > 1$ for some $y \in \mathcal{Y}$, and $f_i \not\equiv f_j$, for some i, j. We have

$$\pi_Y(y|\boldsymbol{\alpha}) = T(\boldsymbol{\alpha}) \sum_{\omega \in \Omega(y)} \frac{\prod_{i=0}^{K} f_i(\omega)^{\alpha_i}}{|\phi'(\omega)|} \quad \text{and} \tag{9}$$

$$\pi'_Y(y|\boldsymbol{\alpha}) = T'(\boldsymbol{\alpha}) \prod_{i=0}^{K} g_i(y)^{\alpha_i} = T'(\boldsymbol{\alpha}) \prod_{i=0}^{K} \left[\sum_{\omega \in \Omega(y)} \frac{f_i(\omega)}{|\phi'(\omega)|} \right]^{\alpha_i}, \tag{10}$$

with

$$T(\boldsymbol{\alpha}) = \int_{\mathcal{Y}} \sum_{\omega \in \Omega(y)} \prod_{i=0}^{K} \left(\frac{f_i(\omega)}{|\phi'(\omega)|} \right)^{\alpha_i} dy = \int_{\mathcal{X}} \prod_{i=0}^{K} f_i(x)^{\alpha_i} dx \quad \text{and} \quad T(\boldsymbol{\alpha})' = \int_{\mathcal{Y}} \prod_{i=0}^{K} \left[\sum_{\omega \in \Omega(y)} \frac{f_i(\omega)}{|\phi'(\omega)|} \right]^{\alpha_i} dy.$$

Now let us construct an explicit example. Let X be a quantity of interest and let $\boldsymbol{F}_X$ be such that each f_i is a Gaussian density with zero mean[2] and variance v_i, i.e.,

$$f_i(x) = \frac{1}{\sqrt{2\pi v_i}} \exp\left(-\frac{x^2}{2v_i} \right)$$

Now consider $\phi(x) = x^2$ and $Y = \phi(X)$. Clearly, $\Omega(y) = \{\omega_0, \omega_1\} = \{-\sqrt{y}, \sqrt{y}\}$ and hence

$$g_i(y) = \frac{f_i(\omega_0)}{|2\omega_0|} + \frac{f_i(\omega_1)}{|2\omega_1|}, \tag{11}$$

$$= \frac{f_i(\sqrt{y})}{\sqrt{y}} = \frac{1}{\sqrt{2\pi v_i y}} \exp\left(-\frac{y}{2v_i} \right). \tag{12}$$

---

[1]Notice there is no guarantee that $|\Omega(y)| < \infty$.
[2]This is important so we can employ a symmetry argument later on.

where the second line follows by using the symmetry of the Gaussian density around zero. We are now prepared to write

$$\pi_Y(y) = \frac{1}{\sqrt{2\pi v^* y}} \exp\left(-\frac{y}{2v^*}\right). \tag{13}$$

whence $v^* = \left[\sum_{i=0}^{K} \frac{\alpha_i}{v_i}\right]^{-1}$. This stems from the pool of Gaussians being also a Gaussian (see manuscript for derivation). An explicit computation gives π_Y':

$$\pi_Y'(y) \propto \prod_{i=0}^{K} g_i(y)^{\alpha_i}, \tag{14}$$

$$\propto \prod_{i=0}^{K} \left(\frac{1}{\sqrt{2\pi v_i y}}\right)^{\alpha_i} \exp\left(-\frac{\alpha_i y}{2v_i}\right), \tag{15}$$

$$= \frac{1}{\sqrt{2\pi v^* y}} \exp\left(-\frac{y}{2v^*}\right). \tag{16}$$

which establishes $\pi_Y(y) = \pi_Y'(y)$ for any $y \in \mathcal{Y}$. □

See Figure 1 for graphical depiction of the example given above. Notice that we could have chosen any distribution on $(-\infty, \infty)$ symmetric about 0. This result shows that it is indeed possible to obtain the same density even when ϕ is not one-to-one.


Figure 1: **Equivalent densities under a non-invertible transform**. Solid line displays $\pi_Y(Y|\boldsymbol{\alpha})$, obtained by first pooling the distributions on X and then computing the induced distribution on Y. Dashed line displays the logarithmic pooling of individual distributions g_i, $\pi_Y'(Y|\boldsymbol{\alpha})$. For this example $K = 5$, $\boldsymbol{\alpha} = \{0.31, 0.08, 0.39, 0.08, 0.15\}$, $\boldsymbol{v} = \{0.25, 0.75, 1.00, 1.00, 0.50\}$.

An example where the two approaches are NOT the same

The order of operations, i.e., whether one adopts procedure (a) or procedure (b), will likely affect the resulting densities, however. In this section we will discuss an example where both procedures lead to densities in the same family, namely the log-normal family of distributions, which are nonetheless distinct.

Suppose $Z = U/V$ and each expert i elicits $U \sim \text{log-normal}(\mu_{iU}, \sigma_{iU}^2)$, $V \sim \text{log-normal}(\mu_{iV}, \sigma_{iV}^2)$, i.e.

$$f_{iU}(u|\mu_{iU}, \sigma_{iU}^2) = \frac{1}{u\sqrt{2\pi\sigma_{iU}^2}} \exp\left(-\frac{(\ln u - \mu_{iU})^2}{2\sigma_{iU}^2}\right),$$

$$g_{iV}(v|\mu_{iV}, \sigma_{iV}^2) = \frac{1}{v\sqrt{2\pi\sigma_{iV}^2}} \exp\left(-\frac{(\ln v - \mu_{iV})^2}{2\sigma_{iV}^2}\right).$$

Again, let $\mathbf{F}_U = \{f_{1U}(U), f_{2U}(U), \ldots, f_{KU}(U)\}$ and $\mathbf{G}_V = \{g_{1V}(V), g_{2V}(V), \ldots, g_{KV}(V)\}$. First, let us derive $\pi_Z(Z)$ under scheme (a). It is not hard to show that $\pi_U(U|\boldsymbol{\alpha}) := \mathcal{LP}(\mathbf{F}_U, \boldsymbol{\alpha}) = \text{log-normal}(\mu_U^*, v_U^*)$ with

$$\mu_U^* := \frac{\sum_{i=0}^K w_{iU}\mu_{iU}}{\sum_{i=0}^K w_{iU}}, \tag{17}$$

$$v_U^* := \frac{1}{\sum_{i=0}^K w_{iU}}, \tag{18}$$

$$w_{iU} := \frac{\alpha_i}{\sigma_{iU}^2}. \tag{19}$$

See our paper for a proof. Analagously, $\pi_V(V|\boldsymbol{\alpha}) := \mathcal{LP}(\mathbf{G}_V, \boldsymbol{\alpha}) = \text{log-normal}(\mu_V^*, v_V^*)$. Then $\pi_Z(Z|\boldsymbol{\alpha}) = \text{log-normal}(\mu_Z^*, v_Z^*)$, with

$$\mu_Z^* = \mu_U^* - \mu_V^*,$$
$$= \frac{\sum_{i=0}^K w_{iU}\mu_{iU}}{\sum_{i=0}^K w_{iU}} - \frac{\sum_{i=0}^K w_{iV}\mu_{iV}}{\sum_{i=0}^K w_{iV}} \quad \text{and} \tag{20}$$

$$v_Z^* = v_U^* + v_V^*,$$
$$= \frac{1}{\sum_{i=0}^K w_{iU}} + \frac{1}{\sum_{i=0}^K w_{iV}}. \tag{21}$$

Now let us consider case (b). Since $r_{iZ} = \text{log-normal}(\mu_{iU} - \mu_{iV}, \sigma_{iU}^2 + \sigma_{iV}^2)$, we arrive at $\pi_Z'(Z|\boldsymbol{\alpha}) = \text{log-normal}(\mu_Z^{**}, v_Z^{**})$,

$$\mu_Z^{**} := \frac{\sum_{i=0}^K w_{iZ}\mu_{iU}}{\sum_{i=0}^K w_{iZ}} - \frac{\sum_{i=0}^K w_{iZ}\mu_{iV}}{\sum_{i=0}^K w_{iZ}}, \tag{22}$$

$$v_Z^{**} := \frac{1}{\sum_{i=0}^K w_{iZ}}, \tag{23}$$

$$w_{iZ} := \frac{\alpha_i}{\sigma_{iU}^2 + \sigma_{iV}^2}. \tag{24}$$

Clearly, $v_Z^* \leq v_Z^{**}$ and hence $\mu_Z^{**} \leq \mu_Z^* \, \forall \, \boldsymbol{\alpha}$.

Minimising Kullback-Leibler divergence in transformed space

One might argue that procedure (b) makes little sense, given that the set of opinions $\mathbf{F}_X$ concerns only X, i.e, it was not necessarily constructed taking the transformation $\phi(\cdot)$ into account. An example is a situation where experts are asked to provide distributions on the probability p of a particular event. In general, elicitation for $f_i(p)$ will not take into account the induced distribution on the log-odds, $\phi(p) = \log p/(1-p)$. Nevertheless, the decision-maker may wish to assign the weights $\boldsymbol{\alpha}$ in a way that takes $\phi(\cdot)$ into account, e.g., by giving lower weights to experts whose distributions on the log-odds scale are unreasonable.

Figure 2: **Log-normal example**. Solid line displays $\pi_Z(Z|\boldsymbol{\alpha})$, obtained by first pooling the distributions on U and V and then computing the induced distribution on Z. Dashed line displays the logarithmic pooling of individual distributions r_{iZ}, $\pi'_Z(Z|\boldsymbol{\alpha})$. For this example $K = 2$, $\alpha_0 = 0.70$, $\mu_{0U} = 0.80$, $\sigma^2_{0U} = 0.40$, $\mu_{1U} = 0.5$, $\sigma^2_{1U} = 0.05$, $\mu_{0V} = -1.60$, $\sigma^2_{0V} = 0.024$, $\mu_{1V} = -1.25$ and $\sigma^2_{1V} = 0.4$.

This decision process can be made more precise. In a similar spirit to the paper, one can construct $\boldsymbol{\alpha}$ so as to minimise the Kullback-Leibler divergence between each distribution in $\mathbf{F_y^{-1}}$ and a transformation of the distribution obtained by procedure (a), $\pi_Y(y|\boldsymbol{\alpha}) = \pi_\theta(\phi^{-1}(y)|\boldsymbol{\alpha})|J_\phi|$. Let $d_i = \mathrm{KL}(h_i(y)||\pi_Y(y|\boldsymbol{\alpha}))$. We then aim at solving the problem

$$L(\boldsymbol{\alpha}) = \sum_{i=0}^{K} d_i \tag{25}$$

$$\hat{\boldsymbol{\alpha}} := \arg\min L(\boldsymbol{\alpha})$$

This procedure therefore choses weights for each expert by how coherent the prior provided by each expert is with the pool-then-Transform – procedure (a) – prior in the transformed space induced by $\phi(\cdot)$.

Acknowledgements

References

Genest, C., Weerahandi, S., and Zidek, J. V. (1984). Aggregating opinions through logarithmic pooling. *Theory and Decision*, 17(1):61–70.