

Savchuk__beta__example.r

luiz

2019-06-24

```
# Leo Bastos & Luiz Max Carvalho (2019)  
# This example was taken from Savchuk and Martz (1994)
```

```
source("pooling_aux.r")
```

```
## Loading required package: ggplot2
```

```
## Registered S3 methods overwritten by 'ggplot2':
```

```
##   method      from  
## [.quosures    rlang  
## c.quosures    rlang  
## print.quosures rlang
```

```
a0 <- 18.1 ; b0 <- .995  
a1 <- 3.44 ; b1 <- .860  
a2 <- 8.32 ; b2 <- .924  
a3 <- 1.98 ; b3 <- .848
```

```
av <- c(a0, a1, a2, a3)  
bv <- c(b0, b1, b2, b3)  
K <- length(av)
```

```
# Individual entropies  
entropies <- rep(NA, K)  
for(k in 1:K) entropies[k] <- entropy_beta(av[k], bv[k])
```

```
entropies
```

```
## [1] -1.9557768 -0.6301164 -1.3069373 -0.2668240
```

```
## Entropy surface (for a future dominance analysis)
```

```
library(fields)
```

```
## Loading required package: spam
```

```
## Loading required package: dotCall64
```

```
## Loading required package: grid
```

```
## Spam version 2.2-2 (2019-03-07) is loaded.
```

```
## Type 'help( Spam)' or 'demo( spam)' for a short introduction  
## and overview of this package.
```

```
## Help for individual functions is also obtained by adding the  
## suffix '.spam' to the function name, e.g. 'help( chol.spam)'.
```

```
##
```

```
## Attaching package: 'spam'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      backsolve, forwardsolve
```

```

## Loading required package: maps

## See https://github.com/NCAR/Fields for
## an extensive vignette, other supplements and source code

ES <- entropy_surface_beta(av, bv)
export <- TRUE
if(export){
  pdf("../plots/entropy_surface_failureProbExample.pdf")
}
image.plot(ES$as, ES$bs, ES$M,
  xlab = expression(a), ylab = expression(b), horizontal = TRUE,
  cex.lab = 1.5, cex.axis = 1.5, axis.args = list(font = 2),
  legend.cex = 1.5,
  legend.lab = expression(H[pi]), main = "Entropy Beta distribution", font = 2)
if(export){
  dev.off()
}

## pdf
## 2

## Observed data
y <- 9
n <- 10

## Marginal likelihoods

marginal.likelihoods <- rep(NA, K)
for (k in 1:K){ marginal.likelihoods[k] <- ml_beta(yi = y, ni = n, a = av[k], b = bv[k]) }

marginal.likelihoods

## [1] 0.2370065 0.2114937 0.2566844 0.1636389

round( normalised.marginal.likelihoods<- marginal.likelihoods/sum(marginal.likelihoods), 2 )

## [1] 0.27 0.24 0.30 0.19

MLS <- marginal_likelihood_surface_beta(y = y, n = n, av = av, bv = bv)
export <- TRUE
if(export){
  pdf("../plots/marginalLikelihood_surface_failureProbExample.pdf")
}
image.plot(MLS$as, MLS$bs, MLS$M,
  xlab = expression(a), ylab = expression(b), horizontal = TRUE,
  cex.lab = 1.5, cex.axis = 1.5, axis.args = list(font = 2),
  legend.cex = 1.5,
  legend.lab = expression(l(y, n)), main = "Marginal likelihood Beta distribution", font = 2)
if(export){
  dev.off()
}

## pdf
## 2

#####
PaperBeta.tbl <- data.frame(mean.prior = rep(NA, 6), lower.prior = NA,

```

```

        upper.prior = NA, mean.post = NA, lower.post = NA,
        upper.post = NA)
rownames(PaperBeta.tbl) <- c("equal_weights", "maximum_entropy", "minimum_KL",
                             "hierarchical_Dirichlet", "hierarchical_LogisticNormal", "Rufo_2012")

AlphasBeta.tbl <- data.frame(matrix(NA, nrow = 5, ncol = length(av)))
rownames(AlphasBeta.tbl) <- c("maximum_entropy", "minimum_KL",
                              "hierarchical_Dirichlet", "hierarchical_LogisticNormal", "Rufo_2012")
colnames(AlphasBeta.tbl) <- paste("alpha_", 0:(K-1), sep = "")


library(ggplot2)

theta.grid <- seq(0, 1, length.out = 1000)
expert.densities <- vector(K, mode = "list")
for(k in 1:K){
  expert.densities[[k]] <- data.frame(theta = theta.grid,
                                     dens = dbeta(theta.grid, shape1 = av[k], shape2 = bv[k]),
                                     expert = paste("expert_", k-1, sep = ""))
}
expert.densities.df <- do.call(rbind, expert.densities)

expert_priors <- ggplot(expert.densities.df, aes(x = theta, y = dens,
                                                linetype = expert, colour = expert)) +
  geom_line(size = 2) +
  scale_linetype_manual(values = c("twodash", "dotted", "longdash", "solid"))+
  scale_colour_brewer(palette = "Spectral") +
  scale_x_continuous(expression(theta), expand = c(0, 0)) +
  scale_y_continuous(expression(f[i](theta)), expand = c(0, 0)) +
  theme_bw(base_size = 20)

expert_priors

```



```
ggsave(expert_priors, filename = "../plots/expert_densities_Savchuk.pdf")
```

```
## Saving 6.5 x 4.5 in image
```

```
##### Equal weights
```

```
alphaEqual <- rep(1/K, K)
```

```
ab.Equal.star <- pool_par(alphaEqual, av, bv)
```

```
# Prior
```

```
(PaperBeta.tbl[1, 1:3] <- stat_beta(ab.Equal.star))
```

```
## [1] 0.8977359 0.6435035 0.9979079
```

```
# Posterior
```

```
(PaperBeta.tbl[1, 4:6] <- stat_beta(ab.Equal.star + c(y, n - y)))
```

```
## [1] 0.8989360 0.7327552 0.9877894
```

```
##### Maximum entropy
```

```
N <- 1000 ## could increase to, say, 10000 in order to make sure, but it's fine
```

```
ent.many.startingPoints <- matrix(rnorm(n = (K-1)*N, mean = 0, sd = 100), ncol = K-1, nrow = N)
```

```
many.ents <- lapply(1:N, function(i) {
```

```
  optim(ent.many.startingPoints[i, ], optentbeta_inv, ap = av, bp = bv)
```


```
})
```

```
optimised.ents <- unlist(lapply(many.ents, function(x) x$value))
```

```
hist(optimised.ents)
```

```
abline(v = optimised.ents[which.min(optimised.ents)], lty = 2, lwd = 2)
```

Histogram of optimised.ents



```
alphaMaxEnt.opt <- alpha_01(many.ents[[which.min(optimised.ents)]]$par)
round(alphaMaxEnt.opt, 2)

## [1] 0 0 0 1
( AlphasBeta.tbl[1, ] <- alphaMaxEnt.opt )

## [1] 2.344230e-38 5.814918e-16 9.743765e-22 1.000000e+00
ab.MaxEnt.star <- pool_par(alphaMaxEnt.opt, av, bv)

# Prior
(PaperBeta.tbl[2, 1:3] <- stat_beta(ab.MaxEnt.star))

## [1] 0.7001414 0.1737705 0.9936577

# Posterior
(PaperBeta.tbl[2, 4:6] <- stat_beta(ab.MaxEnt.star + c(y, n - y)))

## [1] 0.8559401 0.6273426 0.9828843


##### Minimum KL

N <- 1000 ## could increase to, say, 10000 in order to make sure, but it's fine
kl.many.startingPoints <- matrix(rnorm(n = (K-1)*N, mean = 0, sd = 100), ncol = K-1, nrow = N)
many.kls <- lapply(1:N, function(i) {
  optim(kl.many.startingPoints[i, ], optklbeta_inv, ap = av, bp = bv, type = "fp")
})
optimised.kls <- unlist(lapply(many.kls, function(x) x$value))

hist(optimised.kls)
```

```
abline(v = optimised.kls[which.min(optimised.kls)], lty = 2, lwd = 2)
```

Histogram of optimised.kls



```
alphaKL.opt <- alpha_01(many.kls[[which.min(optimised.kls)]]$par)
```

```
(AlphasBeta.tbl[2, ] <- alphaKL.opt)
```

```
## [1] 0.03987021 0.96012979 0.00000000 0.00000000
```

```
ab.KL.star <- pool_par(alphaKL.opt, av, bv)
```

```
# Prior
```

```
(PaperBeta.tbl[3, 1:3] <- stat_beta(ab.KL.star))
```

```
## [1] 0.8230258 0.4242832 0.9966280
```

```
# Posterior
```

```
(PaperBeta.tbl[3, 4:6] <- stat_beta(ab.KL.star + c(y, n-y)))
```

```
## [1] 0.8747215 0.6722137 0.9851126
```

```
##### Hierarchical priors
```

```
require("LearnBayes")
```

```
## Loading required package: LearnBayes
```

```
M <- 100000
```

```
X <- c(1, 1, 1, 1)/10
```

```
alpha.MC.dirichlet <- rdirichlet(M, X)
```

```
alpha.MC.logisticNormal <- rlogisticnorm(N = M,
                                          m = digamma(X)-digamma(X[K]),
                                          Sigma = constructSigma(X))
```

```

apply(alpha.MC.dirichlet, 2, mean)

## [1] 0.2491598 0.2500834 0.2493180 0.2514388
apply(alpha.MC.logisticNormal, 2, mean)

## [1] 0.2492172 0.2512065 0.2500596 0.2495167
apply(alpha.MC.dirichlet, 2, sd)

## [1] 0.3653566 0.3655947 0.3654672 0.3667070
apply(alpha.MC.logisticNormal, 2, sd)

## [1] 0.4019402 0.4034857 0.4026886 0.4023616
beta.par.dirichlet <- alpha.MC.dirichlet %*% cbind(av, bv)
beta.par.logisticNormal <- alpha.MC.logisticNormal %*% cbind(av, bv)

theta.par.dirichlet <- apply(beta.par.dirichlet, 1, function(x) rbeta(1, x[1], x[2]))
theta.par.logisticNormal <- apply(beta.par.logisticNormal, 1, function(x) rbeta(1, x[1], x[2]))
# Prior
PaperBeta.tbl[4, 1] <- mean(theta.par.dirichlet)
PaperBeta.tbl[4, 2:3] <- quantile(theta.par.dirichlet, c(.025, .975))

PaperBeta.tbl[5, 1] <- mean(theta.par.logisticNormal)
PaperBeta.tbl[5, 2:3] <- quantile(theta.par.logisticNormal, c(.025, .975))

##### Hierarchical posteriors

library(rstan)

## Loading required package: StanHeaders
## rstan (Version 2.19.1, GitRev: 2e1f913d3ca3)

## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)
rstan_options(auto_write = TRUE)
options(mc.cores = 4)

betadata.stan <- list(Y = y, X = X, N = n, K = K, a = av, b = bv)

## Dirichlet
compiled.dirichlet <- stan_model("stan/posterior_beta_Dirichlet_pooled.stan")
dirichlet.posterior <- sampling(compiled.dirichlet, data = betadata.stan,
                              control = list(adapt_delta = .99, max_treedepth = 15))

## Warning: There were 120 divergent transitions after warmup. Increasing adapt_delta above 0.99 may help.
## http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup

## Warning: Examine the pairs() plot to diagnose sampling problems
check_hmc_diagnostics(dirichlet.posterior)

##
## Divergences:

```

```

## 120 of 4000 iterations ended with a divergence (3%).
## Try increasing 'adapt_delta' to remove the divergences.

##
## Tree depth:
## 0 of 4000 iterations saturated the maximum tree depth of 15.

##
## Energy:
## E-BFMI indicated no pathological behavior.

theta.dirichlet <- extract(dirichlet.posterior, 'theta')$theta
alphas.dirichlet <- extract(dirichlet.posterior, 'alpha')$alpha
post.alpha.cred.dirichlet <- apply(alphas.dirichlet, 2, quantile, probs = c(.025, .975))
betaPars.dirichlet <- extract(dirichlet.posterior, c('astar', 'bstar'))
ab.dirichlet <- unlist( lapply(betaPars.dirichlet, mean) )

(PaperBeta.tbl[4, 4:6] <- mean_ci(theta.dirichlet) )

##          mean          lwr          upr
## 1 0.8932367 0.7043956 0.9874241
(AlphasBeta.tbl[3, ] <- colMeans(alphas.dirichlet))

## [1] 0.2577346 0.2391749 0.2820408 0.2210496
## Logistic normal
compiled.logisticNormal <- stan_model("stan/posterior_beta_logisticNormal_pooled.stan")

betadata.stan$means <- digamma(X/10)-digamma(X[K]/10)
betadata.stan$Sigma <- constructSigma(X/10)

logisticNormal.posterior <- sampling(compiled.logisticNormal, data = betadata.stan,
                                     control = list(adapt_delta = .99, max_treedepth = 15))

## Warning: There were 136 divergent transitions after warmup. Increasing adapt_delta above 0.99 may help.
## http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup

## Warning: There were 1000 transitions after warmup that exceeded the maximum treedepth. Increase max_treedepth.
## http://mc-stan.org/misc/warnings.html#maximum-treedepth-exceeded

## Warning: Examine the pairs() plot to diagnose sampling problems

## Warning in throw_sampler_warnings(nfits): The largest R-hat is 1.56, indicating chains have not mixed.
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#r-hat

## Warning in throw_sampler_warnings(nfits): Bulk Effective Samples Size (ESS) is too low, indicating poor mixing.
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#bulk-ess

## Warning in throw_sampler_warnings(nfits): Tail Effective Samples Size (ESS) is too low, indicating poor mixing.
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#tail-ess
check_hmc_diagnostics(logisticNormal.posterior)

##
## Divergences:

```



```

## 136 of 4000 iterations ended with a divergence (3.4%).
## Try increasing 'adapt_delta' to remove the divergences.

##
## Tree depth:
## 1000 of 4000 iterations saturated the maximum tree depth of 15 (25%).
## Try increasing 'max_treedepth' to avoid saturation.

##
## Energy:
## E-BFMI indicated no pathological behavior.

theta.logisticNormal <- extract(logisticNormal.posterior, 'theta')$theta
alphas.logisticNormal <- extract(logisticNormal.posterior, 'alpha')$alpha
post.alpha.cred.logisticNormal <- apply(alphas.logisticNormal, 2, quantile, probs = c(.025, .975))
betaPars.logisticNormal <- extract(logisticNormal.posterior, c('astar', 'bstar'))
ab.logisticNormal <- unlist( lapply(betaPars.logisticNormal, mean) )

( PaperBeta.tbl[5, 4:6] <- mean_ci(theta.logisticNormal) )

##          mean          lwr          upr
## 1 0.8899158 0.7108938 0.9853078
( AlphasBeta.tbl[4, ] <- colMeans(alphas.logisticNormal) )

## [1] 0.2160789 0.4340450 0.2163794 0.1334966
##### KL prior from Rufo et al 2012

alphas.rufo <- c(0, 0, 0, 1)

(AlphasBeta.tbl[5, ] <- alphas.rufo)

## [1] 0 0 0 1
ab.rufo <- pool_par(alphas.rufo, av, bv)

# Prior
( PaperBeta.tbl[6, 1:3] <- stat_beta(ab.rufo) )

## [1] 0.7001414 0.1737705 0.9936577

# Posterior
( PaperBeta.tbl[6, 4:6] <- stat_beta(ab.rufo + c(y, n - y)) )

## [1] 0.8559401 0.6273426 0.9828843

#### Finally, tables!
round(PaperBeta.tbl, 3)

##
##          mean.prior lower.prior upper.prior mean.post
## equal_weights          0.898          0.644          0.998          0.899
## maximum_entropy          0.700          0.174          0.994          0.856
## minimum_KL              0.823          0.424          0.997          0.875
## hierarchical_Dirichlet    0.856          0.402          0.998          0.893
## hierarchical_LogisticNormal 0.846          0.351          0.998          0.890
## Rufo_2012                0.700          0.174          0.994          0.856
##
##          lower.post upper.post
## equal_weights          0.733          0.988

```

```
## maximum_entropy          0.627      0.983
## minimum_KL                0.672      0.985
## hierarchical_Dirichlet    0.704      0.987
## hierarchical_LogisticNormal 0.711      0.985
## Rufo_2012                0.627      0.983
```

```
round(AlphasBeta.tbl, 3)
```

```
##                alpha_0 alpha_1 alpha_2 alpha_3
## maximum_entropy    0.000  0.000  0.000  1.000
## minimum_KL         0.040  0.960  0.000  0.000
## hierarchical_Dirichlet 0.258  0.239  0.282  0.221
## hierarchical_LogisticNormal 0.216  0.434  0.216  0.133
## Rufo_2012          0.000  0.000  0.000  1.000
```

```
round(PaperBeta.tbl, 2)
```

```
##                mean.prior lower.prior upper.prior mean.post
## equal_weights          0.90         0.64         1.00      0.90
## maximum_entropy        0.70         0.17         0.99      0.86
## minimum_KL             0.82         0.42         1.00      0.87
## hierarchical_Dirichlet  0.86         0.40         1.00      0.89
## hierarchical_LogisticNormal 0.85         0.35         1.00      0.89
## Rufo_2012              0.70         0.17         0.99      0.86
##                lower.post upper.post
## equal_weights          0.73         0.99
## maximum_entropy        0.63         0.98
## minimum_KL             0.67         0.99
## hierarchical_Dirichlet  0.70         0.99
## hierarchical_LogisticNormal 0.71         0.99
## Rufo_2012              0.63         0.98
```

```
round(AlphasBeta.tbl, 2)
```

```
##                alpha_0 alpha_1 alpha_2 alpha_3
## maximum_entropy    0.00     0.00     0.00     1.00
## minimum_KL         0.04     0.96     0.00     0.00
## hierarchical_Dirichlet 0.26     0.24     0.28     0.22
## hierarchical_LogisticNormal 0.22     0.43     0.22     0.13
## Rufo_2012          0.00     0.00     0.00     1.00
```

```
##### Plotting
```

```
posterior_experts <- data.frame(
  alpha = as.numeric(c(AlphasBeta.tbl[1, ], AlphasBeta.tbl[2, ], AlphasBeta.tbl[5, ],
    AlphasBeta.tbl[3, ], AlphasBeta.tbl[4, ])),
  lwr = c(rep(NA, 12), post.alpha.cred.dirichlet[1, ], post.alpha.cred.logisticNormal[1, ]),
  upr = c(rep(NA, 12), post.alpha.cred.dirichlet[2, ], post.alpha.cred.logisticNormal[2, ]),
  expert = rep(paste("expert_", 0:(K-1), sep = ""), 5),
  method = rep(c("maximum_entropy", "minimum_KL", "Rufo_2012", "Dirichlet", "logistic_normal"), each =
)
```

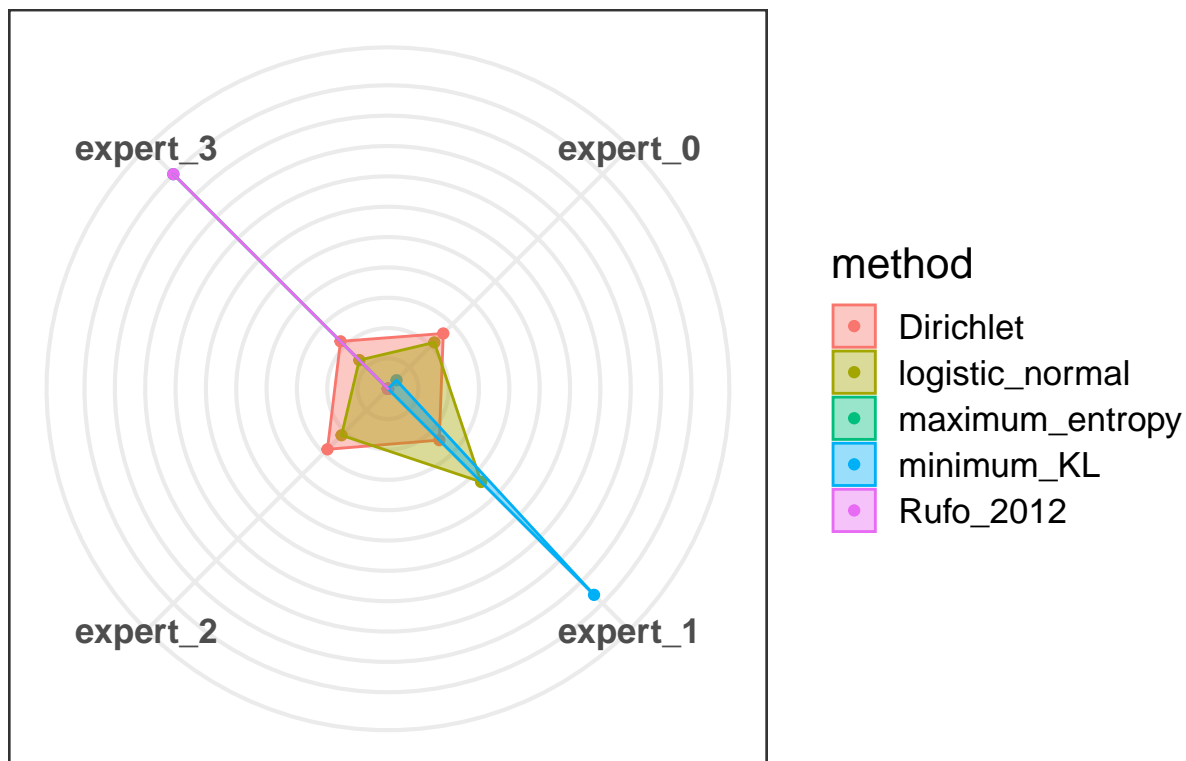
```
####
```

```
radar_alphas <- ggplot(data = posterior_experts,
  aes(x = expert, y = alpha, group = method, colour = method, fill = method)) +
  geom_point() +
  geom_polygon(alpha = 0.4) +
```

```

theme_bw(base_size = 16) +
scale_y_continuous(expand = c(0, 0), limits = c(0, 1),
                    breaks = number_ticks(10)) +
coord_radar() +
theme(axis.title.x = element_blank(),
      axis.ticks.x = element_blank(),
      axis.text.x = element_text(face = "bold"),
      axis.title.y = element_blank(),
      axis.text.y = element_blank(),
      axis.ticks.y = element_blank()
)
radar_alphas

```



```

ggsave(plot = radar_alphas, filename = "../plots/alphas_radar_Savchuk.pdf")

```

```

## Saving 6.5 x 4.5 in image

```

```

#####

```

```

# Now let's look at marginal likelihoods for the pooled priors

```

```

pars <- list(equal_weights = ab.Equal.star,
             maximum_entropy = ab.MaxEnt.star,
             minimum_KL = ab.KL.star,
             hierarchical_Dirichlet = ab.dirichlet ,
             hierarchical_LogisticNormal = ab.logisticNormal,
             Rufo_2012 = ab.rufo)
lapply(pars, function(p) ml_beta(yi = y, ni = n, a = p[1], b = p[2]))

```

```

## $equal_weights

```

```

## [1] 0.2549001

```


```
##
## $maximum_entropy
## [1] 0.1636389
##
## $minimum_KL
## [1] 0.2233347
##
## $hierarchical_Dirichlet
##      astar
## 0.2554928
##
## $hierarchical_LogisticNormal
##      astar
## 0.2536232
##
## $Rufo_2012
## [1] 0.1636389
apply(AlphasBeta.tbl, 1, get_ratio)

##              maximum_entropy              minimum_KL
##              1.719715e+15              2.408138e+01
## hierarchical_Dirichlet hierarchical_LogisticNormal
##              1.094307e+00              2.005944e+00
##              Rufo_2012
##              Inf
J <- length(pars)
posterior.densities.list <- vector(J, mode = "list")
for (j in 1:J){
  posterior.densities.list[[j]] <- data.frame(
    theta = theta.grid,
    dens = dbeta(theta.grid, shape1 = pars[[j]][1], shape2 = pars[[j]][2]),
    method = names(pars)[j]
  )
}

posterior.densities.df <- do.call(rbind, posterior.densities.list)

method_posteriors <- ggplot(posterior.densities.df, aes(x = theta, y = dens,
                                                         linetype = method, colour = method)) +
  geom_line(size = 2) +
  scale_x_continuous(expression(theta), expand = c(0, 0)) +
  scale_y_continuous(expression(pi(theta)), expand = c(0, 0)) +
  geom_vline(xintercept = y/n, linetype = "dashed", size = 1.2) +
  theme_bw(base_size = 16)

method_posteriors
```



```
ggsave(method_posteriors, filename = "../plots/method_posterior_densities_Savchuk.pdf")
```

```
## Saving 6.5 x 4.5 in image
```